

ZILONG (RYAN) WANG

Ph.D., Computer Science, University of California San Diego

zlwang@ucsd.edu

zlwang-cs.github.io

RESEARCH FOCUS

- **Leading research on long CoT reasoning models:** Advancing LLM reasoning capabilities through supervised fine-tuning and reinforcement learning co-design with strategic data selection methodologies.
- **Key contributor to Amazon Shopping Agent:** Specialized in agentic data curation and reward system architecture for large-scale agentic reinforcement learning training pipelines.
- **Secure Code Generation:** Incorporating diverse and accurate reward feedback from static analysis, verification, and additional program analysis tools into the RLVR framework, enabling LLMs to generate functionally correct and security-aware code.

EDUCATION

Doctor of Philosophy in Computer Science

October 2020 – January 2025

University of California, San Diego

- Advisor: Professor Jingbo Shang
- Research Focus: Large language model agents, secure code generation, reinforcement learning

Bachelor of Science in Computer Science

September 2016 – June 2020

Peking University, Beijing, China

- Graduated Summa Cum Laude, Outstanding Graduate Recognition (Top 5%)

PROFESSIONAL EXPERIENCE

Applied Scientist

Present

Amazon, Rufus Team, Palo Alto, CA

- Leading research on long Chain-of-Thought reasoning models for complex problem solving
- Developing large-scale shopping agent systems with RL frameworks and internal search engines
- Designing and implementing agentic data curation pipelines to enhance agent performance and reliability

Research Scientist Intern

June 2024 – December 2024

Amazon, Nile Foundation Team, Palo Alto, CA

- Invented Reward Rising Optimization technique to increase RL data quality in multi-step agentic tasks
- Developed mechanisms for enforcing monotonically increasing reward signals during agent planning processes

Student Researcher

November 2023 – May 2024

Google Cloud AI Research & Google DeepMind

- Developed Speculative RAG framework using specialist models for parallel document-based draft generation
- Designed verification systems using generalist models, achieving 50.83% latency reduction across benchmarks
- Framework integrated into Google Cloud production systems for enterprise deployment

Student Researcher

April 2023 – September 2023

Google Cloud AI Research, Sunnyvale, CA

- Developed Chain-of-Table framework integrating tabular data explicitly within reasoning chains
- Implemented iterative table operations achieving state-of-the-art performance on WikiTQ, FeTaQA, and TabFact benchmarks
- Research deployed to Google Cloud production infrastructure serving enterprise customers

Student Researcher

June 2022 – September 2022

Google DeepMind, Mountain View, CA

- Created VRDU benchmark for multimodal document understanding across forms, receipts, and invoices
- Designed comprehensive evaluation metrics for spatial understanding and hierarchical information extraction

PUBLICATIONS

1. Learning to Optimize Multi-Objective Alignment Through Dynamic Reward Weighting

Yining Lu, Zilong (Ryan) Wang, Shiyang Li, Xin Liu, Changlong Yu, Qingyu Yin, Zhan Shi, Zixuan Zhang, Meng Jiang

Under review at **TACL**

2. ReaL: Training Language Models to Generate Quality Code with Program Analysis Feedback

Zilong (Ryan) Wang, Feng Yao*, Liyuan Liu, Junxia Cui, Li Zhong, Xiaohan Fu, Haohui Mai, Vish Krishnan, Jianfeng Gao, Jingbo Shang (*co-first authorship)*

Conference on Neural Information Processing Systems (**NeurIPS**) 2025

3. RRO: LLM Agent Optimization Through Rising Reward Trajectories

Zilong (Ryan) Wang, Jingfeng Yang, Sreyashi Nag, Samarth Varshney, Xianfeng Tang, Haoming Jiang, Jingbo Shang, Sheikh Muhammad Sarwar

Conference on Language Modeling (**COLM**), 2025

4. Speculative RAG: Enhancing Retrieval Augmented Generation through Drafting

Zilong (Ryan) Wang, Zifeng Wang, Long Le, Huaixiu Steven Zheng, Swaroop Mishra, Vincent Perot, Yuwei Zhang, Anush Mattapalli, Ankur Taly, Jingbo Shang, Chen-Yu Lee, Tomas Pfister

International Conference on Learning Representations (**ICLR**), 2025

5. TableRAG: Million-Token Table Understanding with Language Models

Si-An Chen, Lesly Miculicich, Julian Martin Eisenschlos, Zifeng Wang, Zilong (Ryan) Wang, Yanfei Chen, Yasuhisa Fujii, Hsuan-Tien Lin, Chen-Yu Lee, Tomas Pfister

Conference on Neural Information Processing Systems (**NeurIPS**), 2024

6. Chain-of-Table: Evolving Tables in the Reasoning Chain for Table Understanding

Zilong (Ryan) Wang, Hao Zhang, Chun-Liang Li, Julian Martin Eisenschlos, Vincent Perot, Zifeng Wang, Lesly Miculicich, Yasuhisa Fujii, Jingbo Shang, Chen-Yu Lee, Tomas Pfister

International Conference on Learning Representations (**ICLR**), 2024

7. Debug like a human: A Large Language Model Debugger via Verifying Runtime Execution Step-by-step

Li Zhong, Zilong (Ryan) Wang, Jingbo Shang

Findings of the Association for Computational Linguistics (**ACL**), 2024

8. Answer is All You Need: Instruction-following Text Embedding via Answering the Question

Letian Peng, Zilong (Ryan) Wang, Feng Yao, Zihan Wang, Jingbo Shang

Annual Meeting of the Association for Computational Linguistics (**ACL**), 2024

9. LMDX: Language Model-based Document Information Extraction and Localization

Vincent Perot, Kai Kang, Florian Luisier, Guolong Su, Xiaoyu Sun, Ramya Sree Boppana, Zilong (Ryan) Wang, Jiaqi Mu, Hao Zhang, Nan Hua

Findings of the Association for Computational Linguistics (**ACL**), 2024

10. Can ChatGPT replace StackOverflow? A Study on Robustness and Reliability of Large Language Model Code Generation

Li Zhong, Zilong (Ryan) Wang

The 38th Annual **AAAI** Conference on Artificial Intelligence (**AAAI**), 2023

11. Towards Zero-shot Relation Extraction in Web Mining: A Multimodal Approach with Relative XML Path

Zilong (Ryan) Wang, Jingbo Shang

Findings of the Association for Computational Linguistics (**EMNLP**), 2023

12. VRDU: A Benchmark for Visually-rich Document Understanding

Zilong (Ryan) Wang, Yichao Zhou, Wei Wei, Chen-Yu Lee, Sandeep Tata

Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (**KDD**), 2023

13. MGDDoc: Pre-training with Multi-granular Hierarchy for Document Image Understanding

Zilong (Ryan) Wang, Jiuxiang Gu, Chris Tensmeyer, Nikolaos Barmpalios, Ani Nenkova, Tong Sun, Jingbo Shang, Vlad I. Morariu

Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2022

14. Formulating Few-shot Fine-tuning Towards Language Model Pre-training: A Pilot Study on Named Entity Recognition

Zihan Wang, Kewen Zhao, Zilong (Ryan) Wang, Jingbo Shang

Findings of the Association for Computational Linguistics (EMNLP), 2022

15. Towards Few-shot Entity Recognition in Document Images: A Label-aware Sequence-to-Sequence Framework

Zilong (Ryan) Wang, Jingbo Shang

Findings of the Association for Computational Linguistics (ACL), 2022

16. LayoutReader: Pre-training of Text and Layout for Reading Order Detection

Zilong (Ryan) Wang, Yiheng Xu, Lei Cui, Jingbo Shang, Furu Wei

Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2021

17. DocStruct: A Multimodal Method to Extract Hierarchy Structure in Document for General Form Understanding

Zilong (Ryan) Wang, Mingjie Zhan, Xuebo Liu, Ding Liang

Findings of the Association for Computational Linguistics (EMNLP), 2020

18. Exploring Semantic Capacity of Terms

Jie Huang, Zilong (Ryan) Wang*, Kevin Chang, Wen-mei Hwu, JinJun Xiong (*co-first authorship)*

Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2020

HONORS & SERVICE

Academic Honors: UCSD Powell Fellowship (2020-2021), UCSD Jacob School of Engineering Fellowship (2020-2021)

Professional Service: Reviewer for ICLR, NeurIPS, ICML, ACL, EMNLP, NAACL, CVPR, COLM, COLING